
GLOBAL AND MULTI-OBJECTIVE OPTIMIZATION

Luca Manzoni

ESTIMATION OF DISTRIBUTION ALGORITHMS

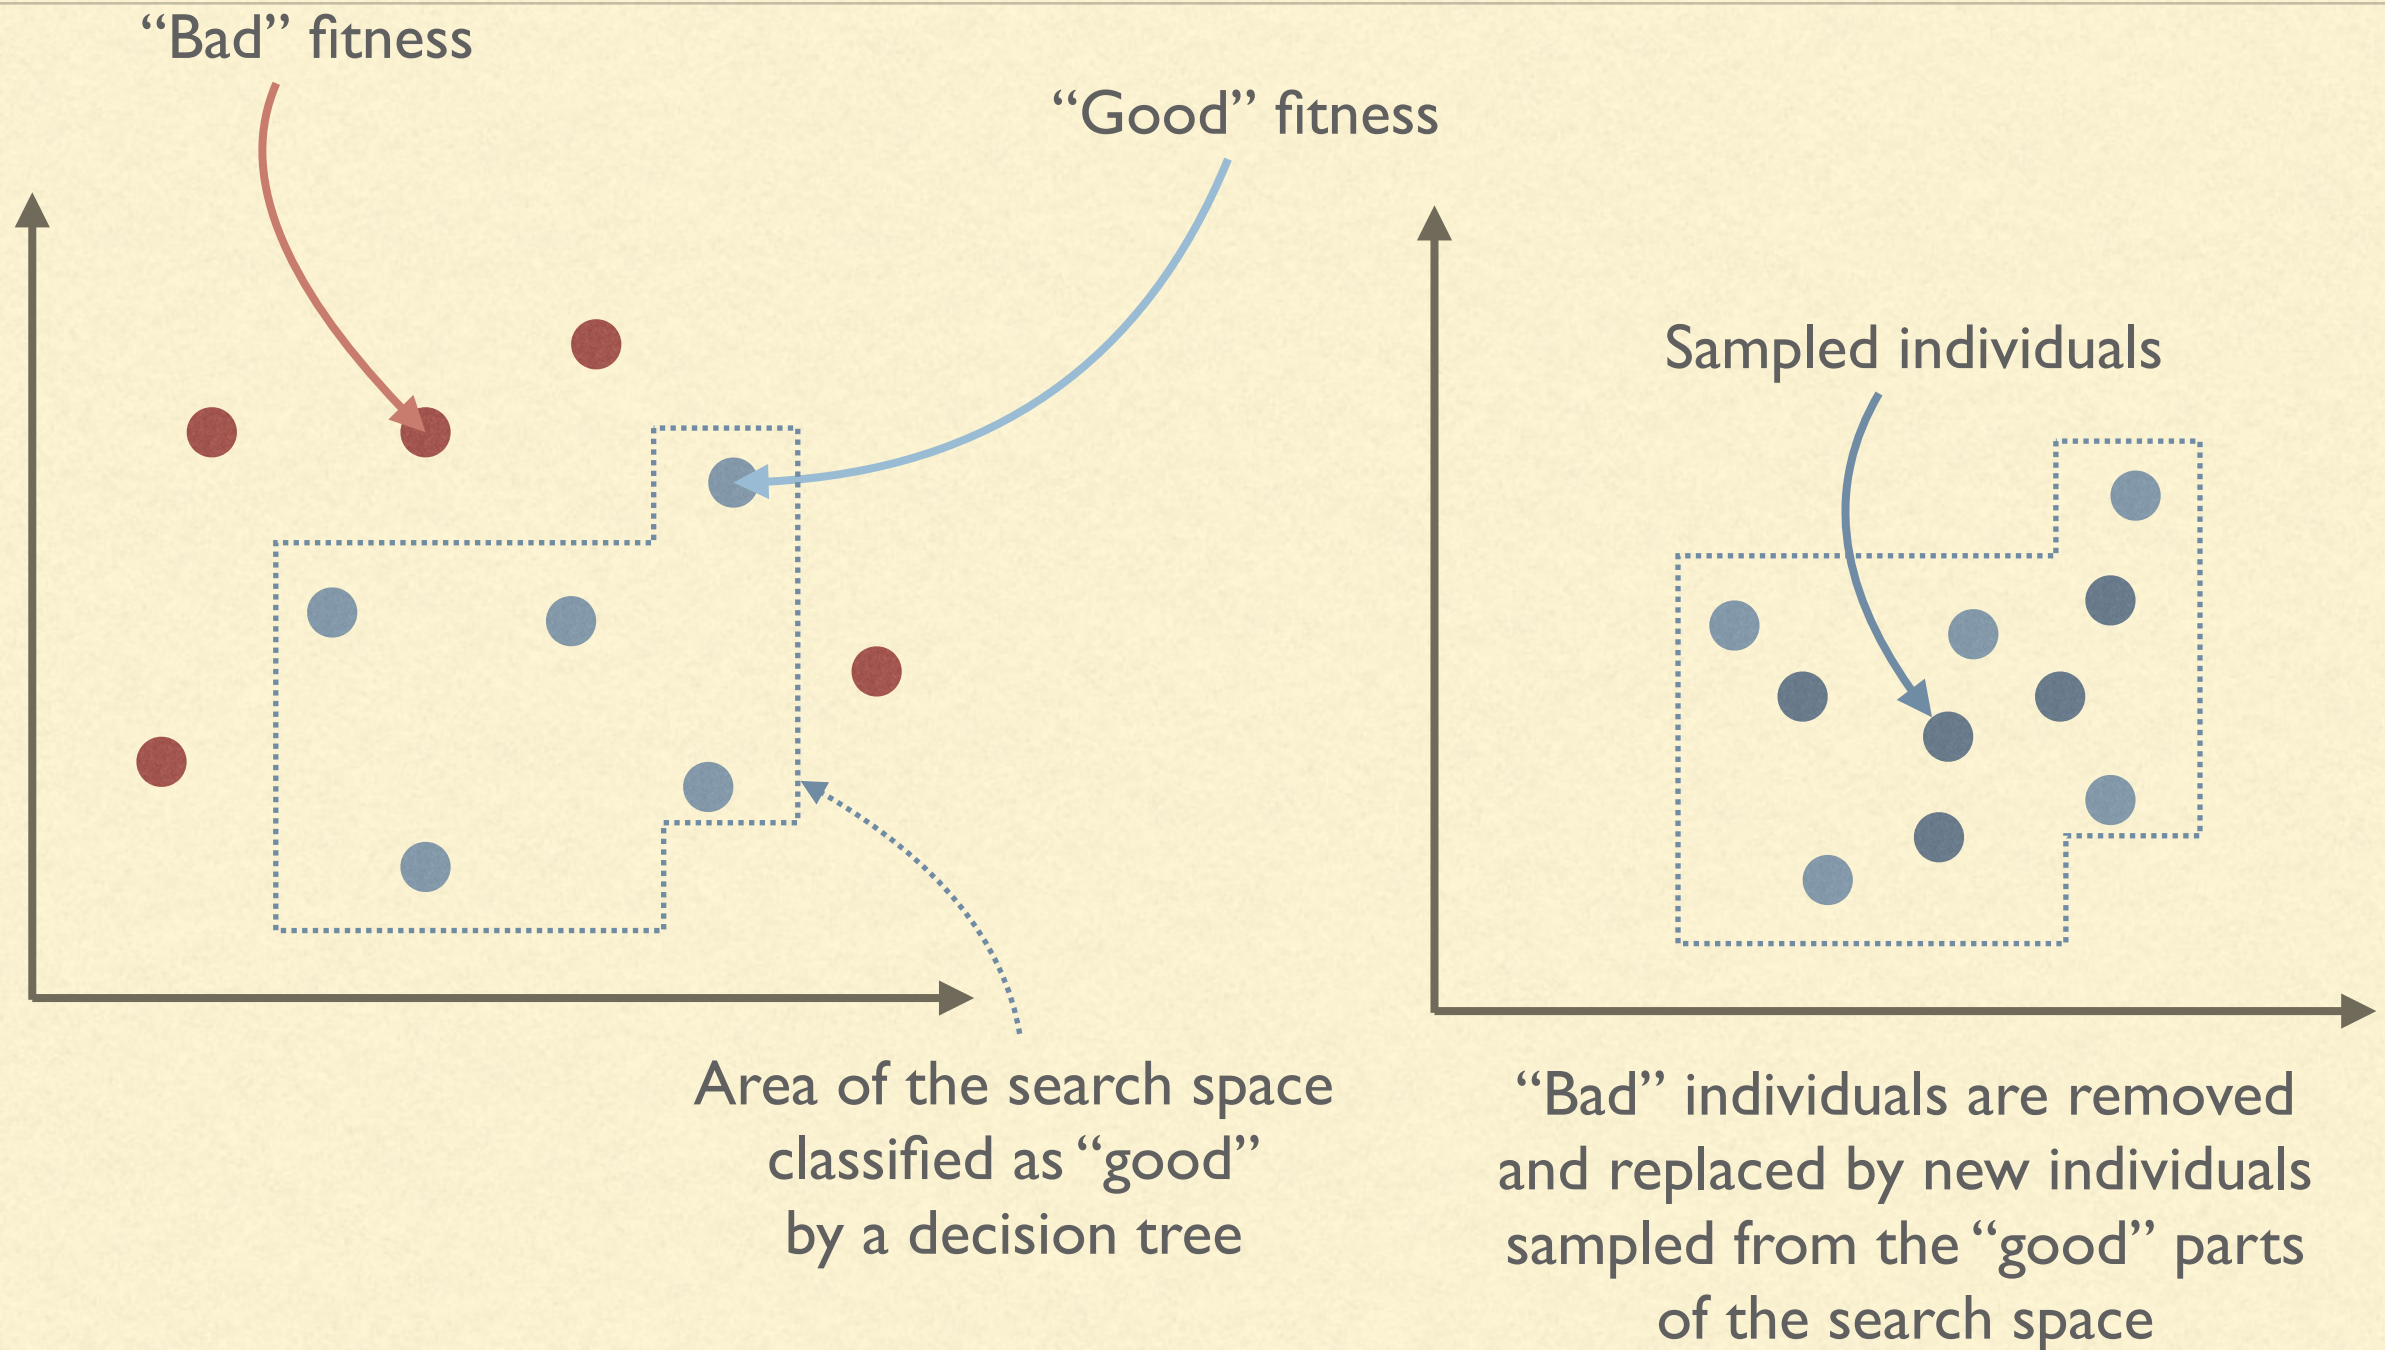
IMPLICIT AND EXPLICIT MODELS

- Selection, recombination (crossover), and mutation are used to sample the search space
 - We may decide to build an explicit model of the space and use it to sample new solutions
 - This is an indirect way of performing selection, recombination, and crossover by using an explicit model
 - The new solutions will then be used to update our model (repeat as many time as desired)
-

MODEL FITTING BY CLASSIFICATION

- We want to sample new solutions in the “good” areas of the search space
 - The idea is to use a **classifier** (e.g., decision trees, SVM, etc) to classify a region of the search space as “good” or “bad” based on the current population
 - We then use the information provided by the classifier to generate a new population
 - The most famous approach are the **Learnable Evolution Models (LEM)**
-

CLASSIFICATION



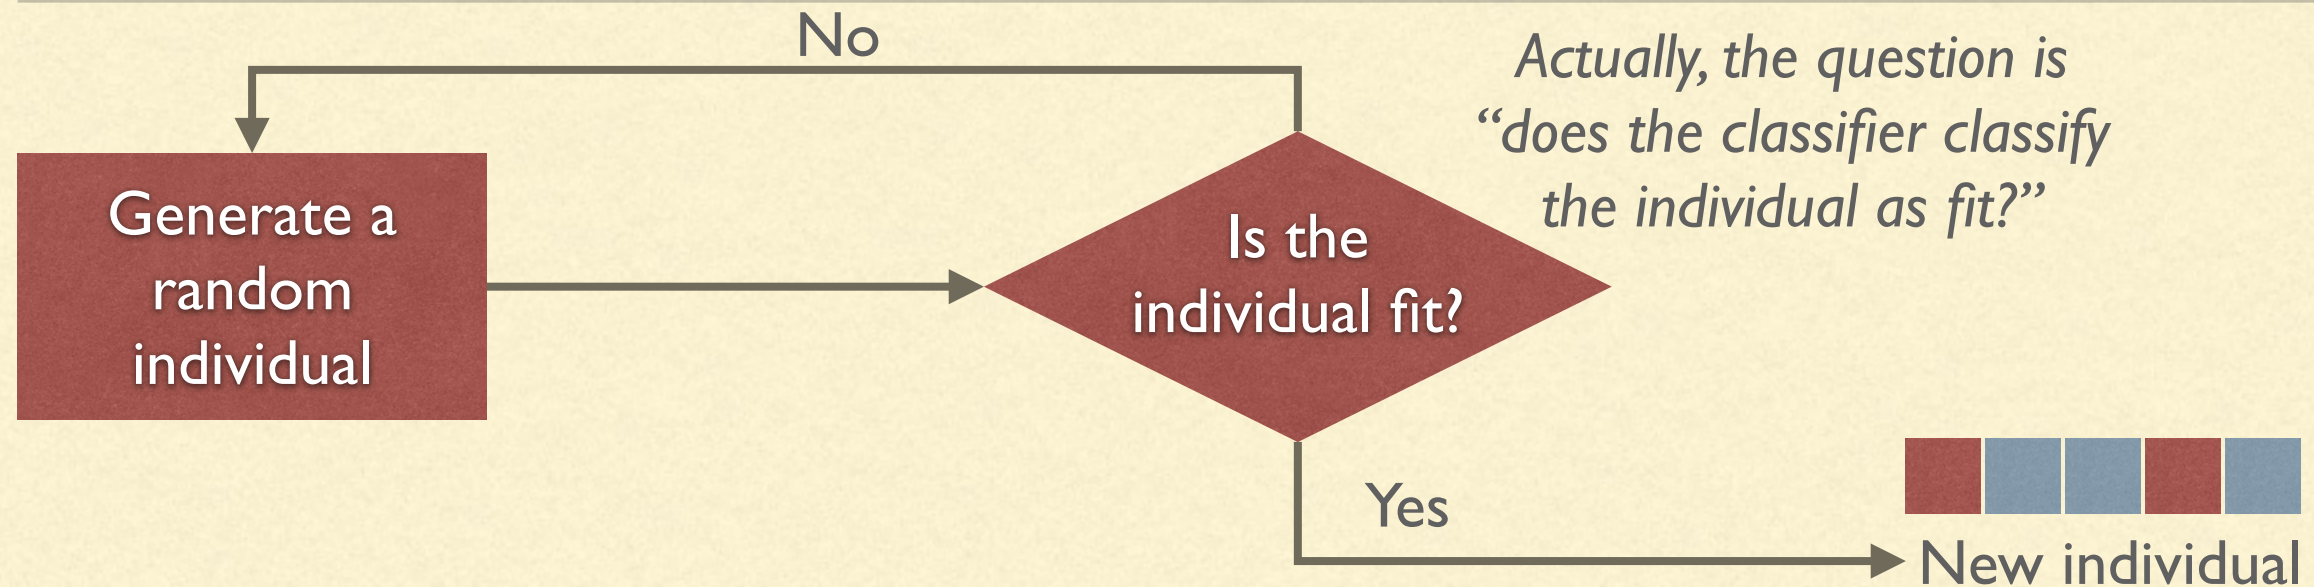
LEARNABLE EVOLUTION MODEL

- Perform some evolutionary steps
 - Divide the population in “fit” and “unfit” individuals
 - Train a classifier (e.g., SVM, Neural Networks, Decision Trees) to distinguish between fit and unfit individuals
 - Remove the unfit individuals and replace them with individuals classified as “fit”
 - Go to the first step unless some termination criteria has been met
-

GENERATIVE VS DISCRIMINATIVE MODELS

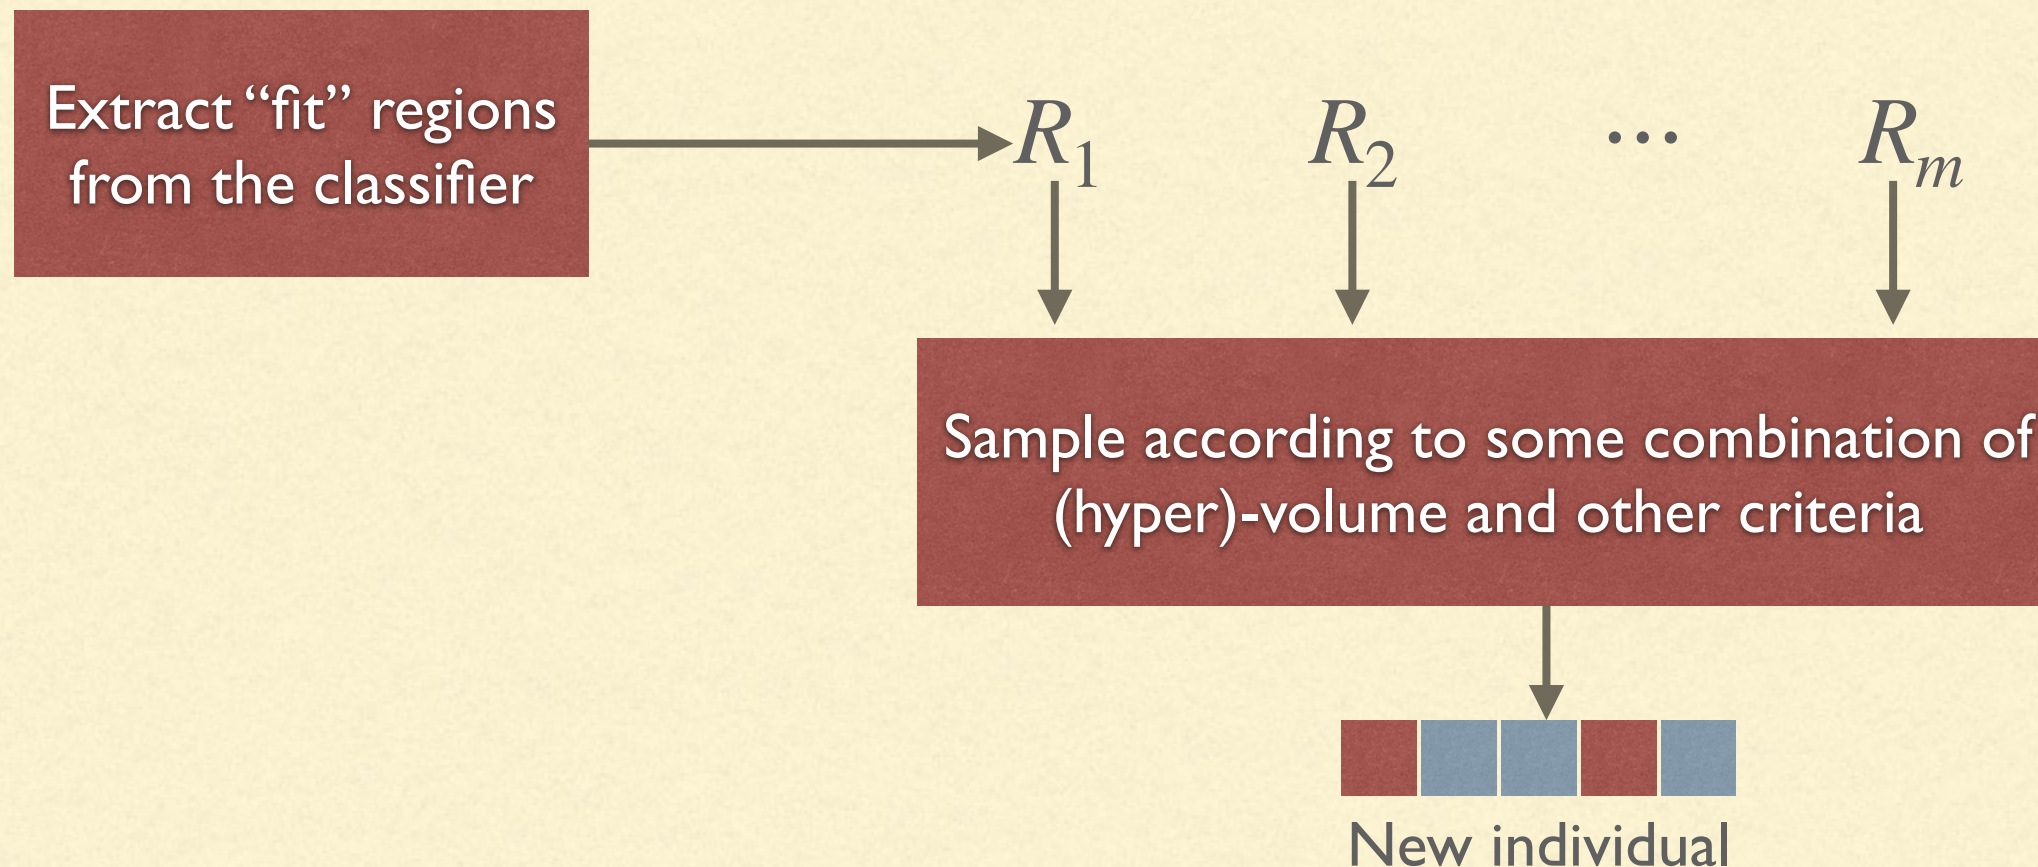
- We can divide the models in two classes:
 - **Generative.** They can be used directly to generate new individuals
 - **Discriminative.** Given an individual they can discriminate if the individual is “good” or “bad”
 - Classifier are generally discriminative models, hence we usually employ them with *rejection sampling*.
-

REJECTION SAMPLING



- We can impose a limit on the number of tries
- After a certain amount of time we have the problem of keeping the number of tries under a reasonable amount
- Most of the search space might be classified as “unfit”, so this sampling becomes too expensive
- We can easily extend it to deal with classifications that are not binary (weighted rejection sampling)

REGION-BASED SAMPLING



- Limited only by our ability to extract fit regions from the classifier
- E.g., perfectly possible with decision trees

WHAT ARE EDA?

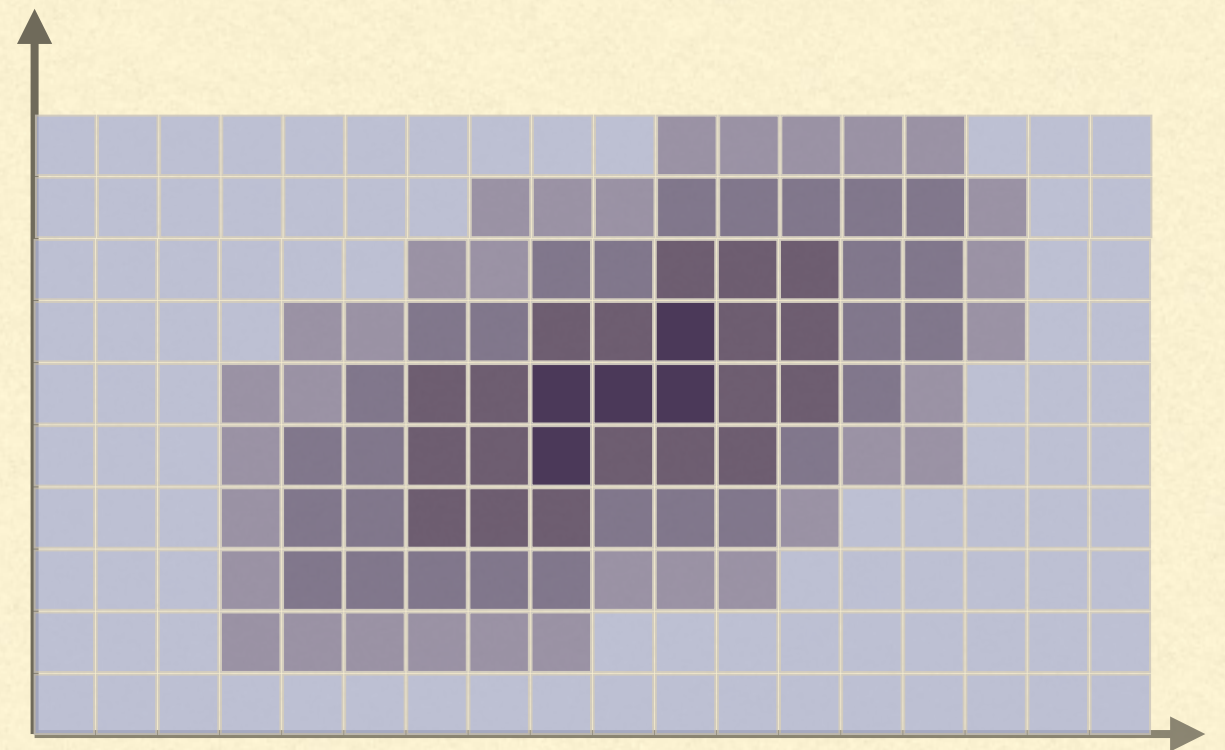
- Sometimes called “probabilistic model-building genetic algorithm” (PMBGA)
 - Instead of using an implicit model or a classifier, EDA use directly a probability distribution to sample the new solutions
-

HISTOGRAMS

The space can be partitioned into (hyper)cubes and we can compute the average fitness of the samples in each of them

If we split each of the n dimensions in d intervals we have d^n hypercubes

Even for $d = 2$ and a small number of dimensions this is unfeasible




A more compact way of representing the fitness distribution should be used

MULTIPLE GAUSSIAN DISTRIBUTIONS

Instead of representing the distribution as an histogram, we can represent it as the sum of a (fixed) number b of Gaussian distributions in n dimensions

More in general, we can use a collection of “known” distributions and sum them




Each Gaussian distribution in n dimensions is determined by:

- The mean vector $\vec{\mu}$ of length n
- The covariance matrix Σ of size n^2


Hence, to represent the distribution as the sum of b Gaussians, we need $b(n + n^2)$ space

MARGINAL DISTRIBUTIONS



Each (one-dimensional) marginal distribution of b Gaussians requires only a mean and a variance, hence the total space used is $2bn$

FINITE DISCRETE SPACES



For Boolean spaces we can represent the marginal distributions as a n -dimensional vector with values in $[0,1]$

UNIVARIATE EDA

- One standard approach with EDA is to marginalise everything to have only distributions in one dimension
 - We are actually assuming that the value of each gene can be determined independently from the values of the others genes
 - We will see two univariate EDA algorithms:
 - Population-based incremental learning (PBIL)
 - Compact Genetic Algorithm (cGA)
-

POPULATION-BASED INCREMENTAL LEARNING

- For dimension n the algorithm starts with n marginal distributions (initially uniform) D_1, \dots, D_n
 - New individuals are sampled picking one gene from each distribution
 - Update the marginal distribution according to the fitness of the sampled individuals
 - Repeat as needed
-

PBIL: UPDATE

- Keep the b fittest individuals in the newly sampled population (i.e., truncated selection)
 - For each gene j , let N_j be the distribution of values of gene j in the b fittest individuals
 - Update $D_j \leftarrow (1 - \alpha)D_j + \alpha N_j$ with $\alpha \in [0,1]$
 - α allows to change the distribution *gradually*
-

PBIL: EXTENSIONS

- PBIL can be extended to work in continuous spaces (e.g., \mathbb{R}^n)
 - One approach is to discrete each marginal distribution into k “buckets”
 - Instead of using a discrete distribution we can use a Gaussian for each marginal distribution
 - In that case each distribution is defined by μ_{D_j} and $\sigma_{D_j}^2$ and the updating rules must modify them
-

PBIL: EXTENSIONS

- Let P be the current population *after the truncated selection* and $P_{i,j}$ the value of gene j in individual i

- We can compute the following values on P

$$\mu_{N_j} = \frac{1}{|P|} \sum_{P_i \in P} P_{i,j} \quad \sigma_{N_j}^2 = \frac{1}{|P| - 1} \sum_{P_i \in P} (P_{i,j} - \mu_{N_j})^2$$


- And the update of the mean and variance of each marginal distribution is performed as

$$\mu_{D_j} \leftarrow (1 - \alpha)\mu_{D_j} + \alpha\mu_{N_j} \quad \sigma_{D_j}^2 \leftarrow (1 - \beta)\sigma_{D_j}^2 + \beta\sigma_{N_j}^2$$

COMPACT GENETIC ALGORITHM

- cGA operates *only* over Boolean spaces
 - cGA update the marginal distributions in steps of a fixed size
 - Instead of computing a new probability distribution, the individuals in the population are compared in pairs
 - For each pair P_i, P_k of individuals, if P_i is fitter than P_k and they differ at gene j , then we shift the distribution D_j to produce the value of gene j in P_i more often
-

COMPACT GENETIC ALGORITHM

- Let $\frac{1}{d}$ be our **discretisation** value (think of it as a “learning rate”). This will be the “step” used to change our distributions
 - We identify D_j with the probability of generating a *one* for gene j
 - Generate a population P of individuals in which each gene is sampled according to D_j
 - For each pair P_i, P_k of individuals, if P_i is fitter than P_k and they differ at gene j , then we shift the distribution D_j to produce the value of gene j in P_i more often (more details later)
- 
- Repeat as needed

COMPACT GENETIC ALGORITHM

- For each pair of individuals $P_i, P_k \in P$
 - Let U be the fittest between P_i and P_k and V the other one
 - If $U_j \neq V_j$, $U_j = 0$, and $D_j > 0$
 - $D_j \leftarrow D_j - \frac{1}{d}$ *shift the distribution toward zero*
 - If $U_j \neq V_j$, $U_j = 1$, and $D_j < 1$
 - $D_j \leftarrow D_j + \frac{1}{d}$ *shift the distribution toward one*
-

ISSUES OF UNIVARIATE EDA

- Univariate EDA assume that the distribution of each gene can be found independently from all the other genes (i.e., no linkage between genes)
 - This is usually false (otherwise we would be able to optimise each gene independently)
 - Univariate EDA can get stuck in local optima due to this
 - **Multivariate EDA** allow to model interaction between genes
-

BAYESIAN OPTIMISATION ALGORITHM

- Among the different multivariate EDA, one is the Bayesian Optimisation Algorithm (BOA)
 - Similar to PBIL where, instead of using marginal distributions, a Bayesian network is used to generate the samples and updated at every generation
 - Pelikan, Martin, David E. Goldberg, and Erick Cantú-Paz.
“*BOA: The Bayesian optimization algorithm.*”
Proceedings of the genetic and evolutionary computation conference GECCO-99.Vol. I. 1999.
-